

# Random Forest

Random Forest is an ensemble of Decision Trees. With a few exceptions, a `RandomForestClassifier` has all the hyperparameters of a `DecisionTreeClassifier` (to control how trees are grown), plus all the hyperparameters of a `BaggingClassifier` to control the ensemble itself.

The Random Forest algorithm introduces extra randomness when growing trees; instead of searching for the very best feature when splitting a node, it searches for the best feature among a random subset of features. This results in a greater tree diversity, which (once again) trades a higher bias for a lower variance, generally yielding an overall better model. The following `BaggingClassifier` is roughly equivalent to the previous `RandomForestClassifier`. Run the cell below to visualize a single estimator from a random forest model, using the Iris dataset to classify the data into the appropriate species.

In [57]:

```
from sklearn.datasets import load_iris
iris = load_iris()

# Model (can also use single decision tree)
from sklearn.ensemble import RandomForestClassifier
model = RandomForestClassifier(n_estimators=10)

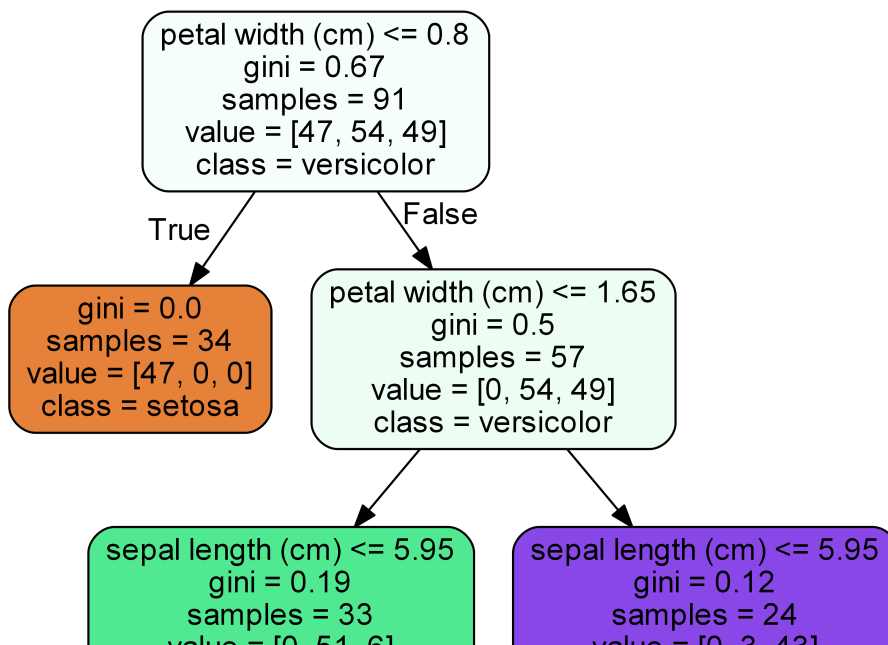
# Train
model.fit(iris.data, iris.target)
# Extract single tree
estimator = model.estimators_[5]

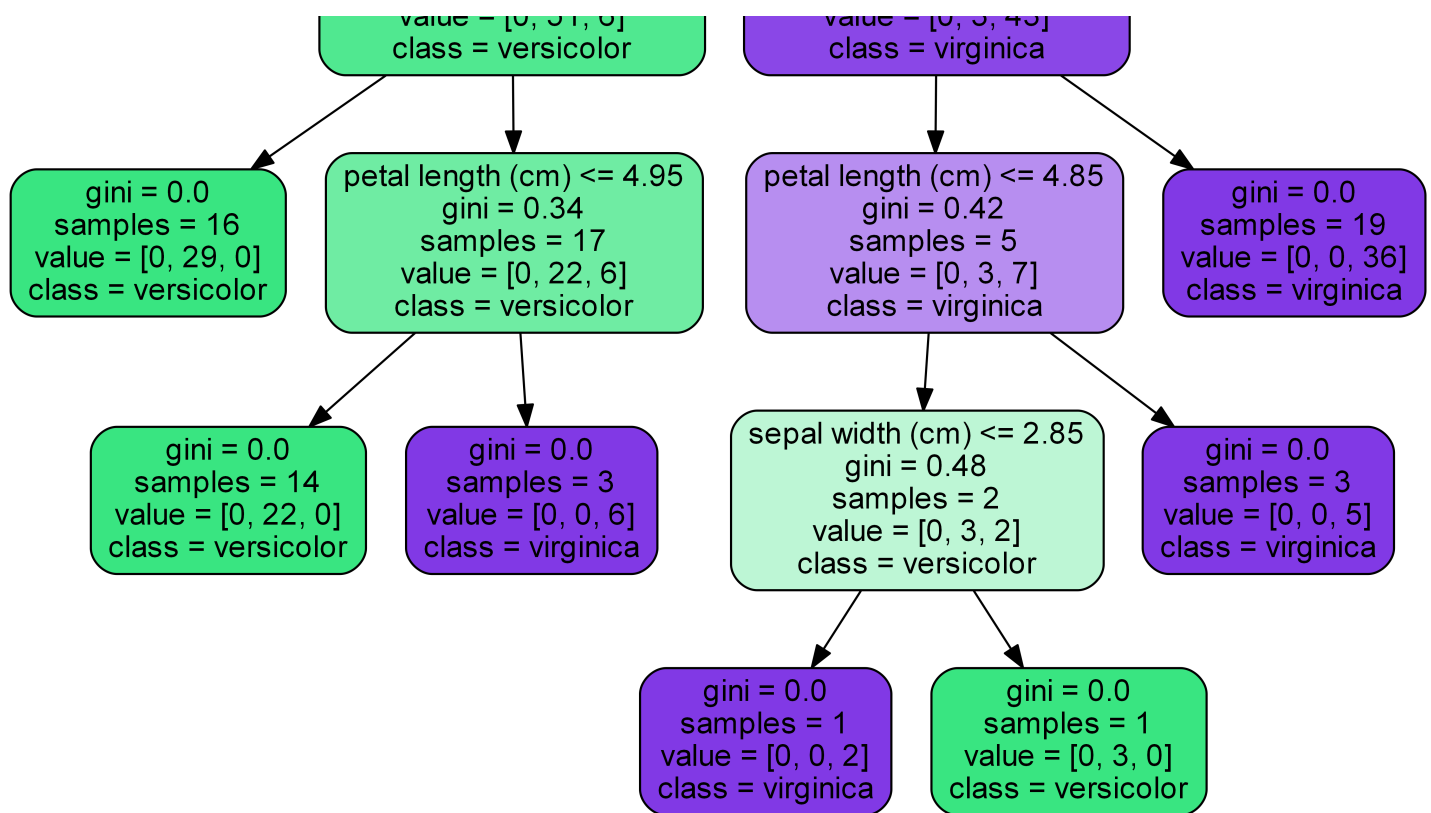
from sklearn.tree import export_graphviz
# Export as dot file
export_graphviz(estimator, out_file='tree.dot',
                 feature_names = iris.feature_names,
                 class_names = iris.target_names,
                 rounded = True, proportion = False,
                 precision = 2, filled = True)

# Convert to png using system command (requires Graphviz)
from subprocess import call
call(['dot', '-Tpng', 'tree.dot', '-o', 'tree.png', '-Gdpi=600'])

# Display in jupyter notebook
from IPython.display import Image
Image(filename = 'tree.png')
```

Out[57]:





Notice how each split separates the data into buckets of similar observations. This is a single tree and a relatively simple classification dataset, but the same method is used in a more complex dataset with greater depth to the trees.

## Coronavirus

Coronavirus disease (COVID-19) is an infectious disease caused by a new virus. The disease causes respiratory illness (like the flu) with symptoms such as a cough, fever, and in more severe cases, difficulty breathing. You can protect yourself by washing your hands frequently, avoiding touching your face, and avoiding close contact (1 meter or 3 feet) with people who are unwell. An outbreak of COVID-19 started in December 2019 and at the time of the creation of this project was continuing to spread throughout the world. Many governments recommended only essential outings to public places and closed most business that do not serve food or sell essential items. An excellent [spatial dashboard](#) built by Johns Hopkins shows the daily confirmed cases by country.

This case study was designed to drive home the important role that data science plays in real-world situations like this pandemic. This case study uses the Random Forest Classifier and a dataset from the South Korean cases of COVID-19 provided on [Kaggle](#) to encourage research on this important topic. The goal of the case study is to build a Random Forest Classifier to predict the 'state' of the patient.

First, please load the needed packages and modules into Python. Next, load the data into a pandas dataframe for ease of use.

In [2]:

```
import os
import pandas as pd
from datetime import datetime, timedelta
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
%matplotlib inline
import plotly.graph_objects as go
from sklearn.experimental import enable_iterative_imputer
from sklearn.impute import IterativeImputer
from sklearn.ensemble import ExtraTreesRegressor
```

In [3]:

```
#url = 'SouthKoreacoronavirusdataset/PatientInfo.csv'
df = pd.read_csv(r'C:\Users\admin\Desktop\SPRINGBOARDFILES\Unit 14\Unit 14.4\1593492590_RandomForest_Covid_Case_Study_06302020\RandomForest Covid Case Study_06302020\PatientInfo.csv')
df.head()
```

Out[3]:

	patient_id	global_num	sex	birth_year	age	country	province	city	disease	infection_case	infection_order
0	1000000001	2.0	male	1964.0	50s	Korea	Seoul	Gangseo-gu	NaN	overseas inflow	1.0
1	1000000002	5.0	male	1987.0	30s	Korea	Seoul	Jungnang-gu	NaN	overseas inflow	1.0
2	1000000003	6.0	male	1964.0	50s	Korea	Seoul	Jongno-gu	NaN	contact with patient	2.0 2
3	1000000004	7.0	male	1991.0	20s	Korea	Seoul	Mapo-gu	NaN	overseas inflow	1.0
4	1000000005	9.0	female	1992.0	20s	Korea	Seoul	Seongbuk-gu	NaN	contact with patient	2.0 1

In [4]:

```
df.shape
```

Out[4]:

(2218, 18)

In [5]:

```
df["disease"].unique()
```

Out[5]:

array([nan, True], dtype=object)

In [6]:

```
#Counts of null values
na_df=pd.DataFrame(df.isnull().sum().sort_values(ascending=False)).reset_index()
na_df.columns = ['VarName', 'NullCount']
na_df[(na_df['NullCount']>0)]
```

Out[6]:

	VarName	NullCount
0	disease	2199
1	deceased_date	2186
2	infection_order	2176
3	symptom_onset_date	2025
4	released_date	1995
5	contact_number	1807
6	infected_by	1749
7	infection_case	1055
8	global_num	904
9	birth_year	454
10	age	261
11	sex	145
12	confirmed_date	141

13	state	88
VarName	NullCount	
14	city	65

In [7]:

```
#counts of response variable values
df.state.value_counts()
```

Out[7]:

```
isolated    1791
released     307
deceased     32
Name: state, dtype: int64
```

**Create a new column named 'n\_age' which is the calculated age based on the birth year column.**

In [8]:

```
df['n_age']=2020- df['birth_year']
df['n_age']
```

Out[8]:

```
0      56.0
1      33.0
2      56.0
3      29.0
4      28.0
...
2213    30.0
2214    22.0
2215    22.0
2216    48.0
2217    46.0
Name: n_age, Length: 2218, dtype: float64
```

## Handle Missing Values

**Print the number of missing values by column.**

In [9]:

```
df.isnull().sum()
```

Out[9]:

```
patient_id      0
global_num     904
sex            145
birth_year     454
age            261
country         0
province        0
city           65
disease        2199
infection_case  1055
infection_order 2176
infected_by     1749
contact_number  1807
symptom_onset_date 2025
confirmed_date   141
released_date   1995
deceased_date   2186
state           88
n_age          454
dtype: int64
```

In [10]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2218 entries, 0 to 2217
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  ---                ---
0   patient_id            2218 non-null   int64
1   global_num            1314 non-null   float64
2   sex                   2073 non-null   object
3   birth_year            1764 non-null   float64
4   age                   1957 non-null   object
5   country               2218 non-null   object
6   province              2218 non-null   object
7   city                  2153 non-null   object
8   disease               19 non-null     object
9   infection_case        1163 non-null   object
10  infection_order        42 non-null     float64
11  infected_by            469 non-null     float64
12  contact_number         411 non-null     float64
13  symptom_onset_date     193 non-null     object
14  confirmed_date         2077 non-null     object
15  released_date          223 non-null     object
16  deceased_date          32 non-null     object
17  state                  2130 non-null     object
18  n_age                  1764 non-null     float64
dtypes: float64(6), int64(1), object(12)
memory usage: 225.3+ KB
```

**Fill the 'disease' missing values with 0 and remap the True values to 1.**

```
In [11]:
```

```
df["disease"] = df["disease"].fillna(0)
```

```
In [12]:
```

```
df["disease"] = df["disease"].replace(True,1)
```

```
In [13]:
```

```
df["disease"].unique()
```

```
Out[13]:
```

```
array([0, 1], dtype=int64)
```

**Fill null values in the following columns with their mean:  
'global\_number','birth\_year','infection\_order','infected\_by'and 'contact\_number'**

```
In [14]:
```

```
df['global_num'].fillna((df['global_num'].mean()), inplace=True)
```

```
In [15]:
```

```
df['birth_year'].fillna((df['birth_year'].mean()), inplace=True)
```

```
In [16]:
```

```
df['infection_order'].fillna((df['infection_order'].mean()), inplace=True)
```

```
In [17]:
```

```
df['infected_by'].fillna((df['infected_by'].mean()), inplace=True)
```

```
In [18]:
```

```
df['contact_number'].fillna((df['contact_number'].mean()), inplace=True)
```

Fill the rest of the missing values with any method.

In [19]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2218 entries, 0 to 2217
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   patient_id            2218 non-null   int64
1   global_num            2218 non-null   float64
2   sex                   2073 non-null   object
3   birth_year            2218 non-null   float64
4   age                   1957 non-null   object
5   country               2218 non-null   object
6   province              2218 non-null   object
7   city                  2153 non-null   object
8   disease               2218 non-null   int64
9   infection_case        1163 non-null   object
10  infection_order        2218 non-null   float64
11  infected_by            2218 non-null   float64
12  contact_number         2218 non-null   float64
13  symptom_onset_date     193 non-null    object
14  confirmed_date         2077 non-null   object
15  released_date          223 non-null    object
16  deceased_date          32 non-null     object
17  state                  2130 non-null   object
18  n_age                  1764 non-null   float64
dtypes: float64(6), int64(2), object(11)
memory usage: 234.0+ KB
```

In [20]:

```
df['sex'].fillna((df['sex'].mode()[0]), inplace=True)
```

In [21]:

```
df['age'].fillna((df['age'].mode()[0]), inplace=True)
```

In [22]:

```
df['city'].fillna((df['city'].mode()[0]), inplace=True)
```

In [23]:

```
df['infection_case'].fillna((df['infection_case'].mode()[0]), inplace=True)
```

In [24]:

```
df['symptom_onset_date'].fillna((df['symptom_onset_date'].mode()[0]), inplace=True)
```

In [25]:

```
df['confirmed_date'].fillna((df['confirmed_date'].mode()[0]), inplace=True)
```

In [26]:

```
df['deceased_date'].fillna((df['deceased_date'].mode()[0]), inplace=True)
```

In [27]:

```
df['state'].fillna((df['state'].mode()[0]), inplace=True)
```

In [28]:

```
df['n_age'].fillna((df['n_age'].mode()[0]), inplace=True)
```

Check for any remaining null values.

In [29]:

```
df.isnull()
```

Out[29]:

	patient_id	global_num	sex	birth_year	age	country	province	city	disease	infection_case	infection_order	infect
0	False	False	False	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False	False	False	False
...	...	...	...	...	...	...	...	...	...	...	...	...
2213	False	False	False	False	False	False	False	False	False	False	False	False
2214	False	False	False	False	False	False	False	False	False	False	False	False
2215	False	False	False	False	False	False	False	False	False	False	False	False
2216	False	False	False	False	False	False	False	False	False	False	False	False
2217	False	False	False	False	False	False	False	False	False	False	False	False

2218 rows x 19 columns



In [30]:

```
df.head()
```

Out[30]:

	patient_id	global_num	sex	birth_year	age	country	province	city	disease	infection_case	infection_order
0	1000000001	2.0	male	1964.0	50s	Korea	Seoul	Gangseo-gu	0	overseas inflow	1.0 2
1	1000000002	5.0	male	1987.0	30s	Korea	Seoul	Junngnang-gu	0	overseas inflow	1.0 2
2	1000000003	6.0	male	1964.0	50s	Korea	Seoul	Jongno-gu	0	contact with patient	2.0 2
3	1000000004	7.0	male	1991.0	20s	Korea	Seoul	Mapo-gu	0	overseas inflow	1.0 2
4	1000000005	9.0	female	1992.0	20s	Korea	Seoul	Seongbuk-gu	0	contact with patient	2.0 1



Remove date columns from the data.

In [31]:

```
df = df.drop(['symptom_onset_date', 'confirmed_date', 'released_date', 'deceased_date'], axis = 1)
```

Review the count of unique values by column.

In [32]:

```
print(df.nunique())
```

patient\_id 2218  
global\_num 1304  
^

```

sex                2
birth_year         97
age                11
country            4
province          17
city              134
disease            2
infection_case     16
infection_order    7
infected_by       207
contact_number     73
state              3
n_age             96
dtype: int64

```

Review the percent of unique values by column.

In [33]:

```
print(df.nunique()/df.shape[0])
```

```

patient_id        1.000000
global_num        0.587917
sex               0.000902
birth_year        0.043733
age               0.004959
country           0.001803
province          0.007665
city              0.060415
disease           0.000902
infection_case    0.007214
infection_order   0.003156
infected_by       0.093327
contact_number    0.032913
state             0.001353
n_age             0.043282
dtype: float64

```

Review the range of values per column.

In [34]:

```
df.describe().T
```

Out[34]:

	count	mean	std	min	25%	50%	75%	max
<b>patient_id</b>	2218.0	4.014678e+09	2.192419e+09	1.000000e+09	1.700000e+09	6.001000e+09	6.004000e+09	7.000000e+09
<b>global_num</b>	2218.0	4.664817e+03	2.211785e+03	1.000000e+00	4.205250e+03	4.664817e+03	5.900250e+03	8.717000e+03
<b>birth_year</b>	2218.0	1.974989e+03	1.731123e+01	1.916000e+03	1.965000e+03	1.974989e+03	1.988000e+03	2.020000e+03
<b>disease</b>	2218.0	8.566276e-03	9.217769e-02	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	1.000000e+00
<b>infection_order</b>	2218.0	2.285714e+00	1.706622e-01	1.000000e+00	2.285714e+00	2.285714e+00	2.285714e+00	6.000000e+00
<b>infected_by</b>	2218.0	2.600789e+09	7.216328e+08	1.000000e+09	2.600789e+09	2.600789e+09	2.600789e+09	6.113000e+09
<b>contact_number</b>	2218.0	2.412895e+01	3.917141e+01	0.000000e+00	2.412895e+01	2.412895e+01	2.412895e+01	1.160000e+03
<b>n_age</b>	2218.0	4.623715e+01	1.747912e+01	0.000000e+00	3.200000e+01	5.100000e+01	5.500000e+01	1.040000e+02

Check for duplicated rows

In [35]:

```

duplicateRowsDF = df[df.duplicated()]
duplicateRowsDF

```



Out [35]:

patient_id	global_num	sex	birth_year	age	country	province	city	disease	infection_case	infection_order	infected_by	contact_number	n_age
------------	------------	-----	------------	-----	---------	----------	------	---------	----------------	-----------------	-------------	----------------	-------

Print the categorical columns and their associated levels.

In [36]:

```
dfo = df.select_dtypes(include=['object'], exclude=['datetime'])
dfo.shape
#get levels for all variables
vn = pd.DataFrame(dfo.nunique()).reset_index()
vn.columns = ['VarName', 'LevelsCount']
vn.sort_values(by='LevelsCount', ascending =False)
vn
```

Out [36]:

	VarName	LevelsCount
0	sex	2
1	age	11
2	country	4
3	province	17
4	city	134
5	infection_case	16
6	state	3

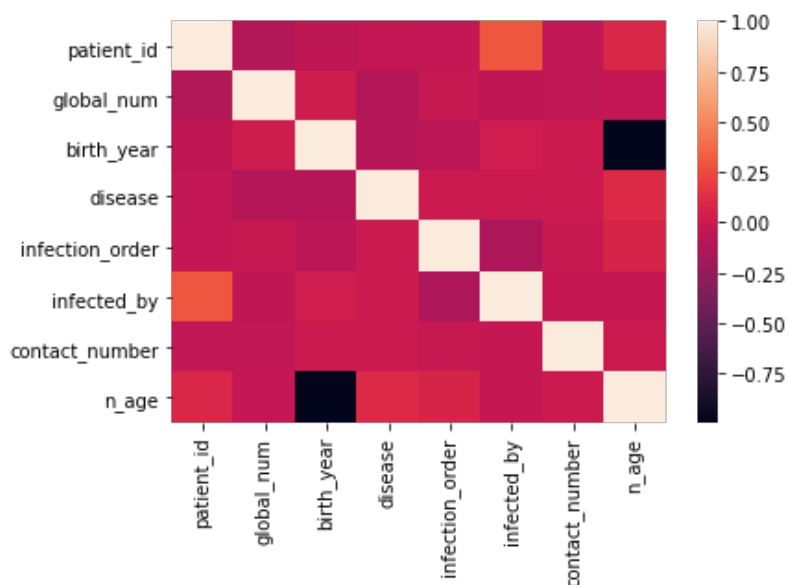
Plot the correlation heat map for the features.

In [37]:

```
sns.heatmap(df.corr())
```

Out [37]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x8ffe2d0>




Plot the boxplots to check for outliers.

In [38]:

```
plt.figure(figsize=(20,10))
df.boxplot()
```

Out[38]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x944e170>



Create dummy features for object type features.

In [39]:

```
features=['sex', 'age', 'country', 'province', 'city', 'infection_case']
dummies=pd.get_dummies(df[features])
merged=pd.concat([df,dummies],axis=1)
final=merged.drop(['sex', 'age', 'country', 'province', 'city', 'infection_case'], axis=1)
df=final
df.head()
```

Out[39]:

	patient_id	global_num	birth_year	disease	infection_order	infected_by	contact_number	state	n_age	sex_female
0	1000000001	2.0	1964.0	0	1.0	2.600789e+09	75.0	released	56.0	0
1	1000000002	5.0	1987.0	0	1.0	2.600789e+09	31.0	released	33.0	0
2	1000000003	6.0	1964.0	0	2.0	2.002000e+09	17.0	released	56.0	0
3	1000000004	7.0	1991.0	0	1.0	2.600789e+09	9.0	released	29.0	0
4	1000000005	9.0	1992.0	0	2.0	1.000000e+09	2.0	released	28.0	1

5 rows x 193 columns



Split the data into test and train subsamples

In [40]:

```
df.columns
```

Out[40]:

```
Index(['patient_id', 'global_num', 'birth_year', 'disease', 'infection_order',
      'infected_by', 'contact_number', 'state', 'n_age', 'sex_female',
      ...,
      'infection_case_Pilgrimage to Israel',
      'infection case River of Grace Community Church',
```

```
'infection_case_Seongdong-gu APT', 'infection_case_Shincheonji Church',  
'infection_case_Suyeong-gu Kindergarten',  
'infection_case_contact with patient', 'infection_case_etc',  
'infection_case_gym facility in Cheonan',  
'infection_case_gym facility in Sejong',  
'infection_case_overseas inflow'],  
dtype='object', length=193)
```

In [45]:

```
from sklearn.model_selection import train_test_split  
  
# dont forget to define your X and y  
X= df.drop(['state'],axis=1)  
y=df['state']  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.2, random_state=1)  
X_train = pd.get_dummies(X_train)  
X_test = pd.get_dummies(X_test)
```

## Scale data to prep for model creation

In [46]:

```
#scale data  
from sklearn import preprocessing  
import numpy as np  
# build scaler based on training data and apply it to test data to then also scale the te  
st data  
scaler = preprocessing.StandardScaler().fit(X_train)  
X_train_scaled=scaler.transform(X_train)  
X_test_scaled=scaler.transform(X_test)
```

In [47]:

```
from sklearn.metrics import precision_recall_curve  
from sklearn.metrics import f1_score  
from sklearn.metrics import auc  
from sklearn.linear_model import LogisticRegression  
from matplotlib import pyplot  
from sklearn.metrics import precision_recall_curve  
from sklearn.metrics import f1_score  
from sklearn.metrics import auc  
from sklearn.linear_model import LogisticRegression  
from sklearn.metrics import classification_report, confusion_matrix, roc_curve, roc_auc_scor  
e  
from sklearn.metrics import accuracy_score, log_loss  
from matplotlib import pyplot
```

## Fit Random Forest Classifier

The fit model shows an overall accuracy of 80% which is great and indicates our model was effectively able to identify the status of a patients in the South Korea dataset.

In [48]:

```
from sklearn.ensemble import RandomForestClassifier  
clf = RandomForestClassifier(n_estimators=300, random_state = 1,n_jobs=-1)  
model_res = clf.fit(X_train_scaled, y_train)  
y_pred = model_res.predict(X_test_scaled)  
y_pred_prob = model_res.predict_proba(X_test_scaled)  
lr_probs = y_pred_prob[:,1]  
ac = accuracy_score(y_test, y_pred)  
  
f1 = f1_score(y_test, y_pred, average='weighted')  
cm = confusion_matrix(y_test, y_pred)  
  
print('Random Forest: Accuracy=%.3f' % (ac))
```

```
print('Random Forest: f1-score=%.3f' % (f1))
```

Random Forest: Accuracy=0.865

Random Forest: f1-score=0.832

## Create Confusion Matrix Plots

Confusion matrices are great ways to review your model performance for a multi-class classification problem. Being able to identify which class the misclassified observations end up in is a great way to determine if you need to build additional features to improve your overall model. In the example below we plot a regular counts confusion matrix as well as a weighted percent confusion matrix. The percent confusion matrix is particularly helpful when you have unbalanced class sizes.

In [49]:

```
class_names=['isolated','released','missing','deceased'] # name of classes
```

In [50]:

```
import itertools
import numpy as np
import matplotlib.pyplot as plt

from sklearn import svm, datasets
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix

def plot_confusion_matrix(cm, classes,
                          normalize=False,
                          title='Confusion matrix',
                          cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

    print(cm)

    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

    fmt = '.2f' if normalize else 'd'
    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, format(cm[i, j], fmt),
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.ylabel('True label')
    plt.xlabel('Predicted label')
    plt.tight_layout()

# Compute confusion matrix
cnf_matrix = confusion_matrix(y_test, y_pred)
np.set_printoptions(precision=2)

# Plot non-normalized confusion matrix
plt.figure()
plot_confusion_matrix(cnf_matrix, classes=class_names,
```

```

        title='Confusion matrix, without normalization')
plt.savefig('figures/RF_cm_multi_class.png')

# Plot normalized confusion matrix
plt.figure()
plot_confusion_matrix(cnf_matrix, classes=class_names, normalize=True,
                      title='Normalized confusion matrix')
plt.savefig('figures/RF_cm_proportion_multi_class.png', bbox_inches="tight")
plt.show()

```

Confusion matrix, without normalization

```

[[ 6  1  0]
 [ 0 363  5]
 [ 0  54 15]]


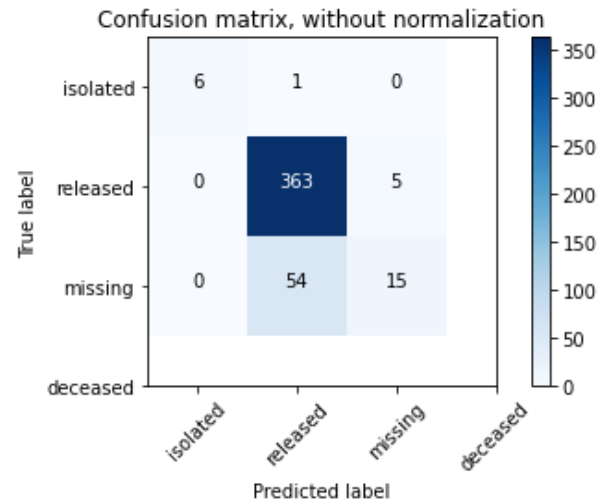
```

Normalized confusion matrix

```

[[0.86 0.14 0.  ]
 [0.   0.99 0.01]
 [0.   0.78 0.22]]

```



## Plot feature importances

The random forest algorithm can be used as a regression or classification model. In either case it tends to be a bit of a black box, where understanding what's happening under the hood can be difficult. Plotting the feature importances is one way that you can gain a perspective on which features are driving the model predictions.

In [51]:

```


feature_importance = clf.feature_importances_
# make importances relative to max importance
feature_importance = 100.0 * (feature_importance / feature_importance.max())[:30]
sorted_idx = np.argsort(feature_importance)[:30]

pos = np.arange(sorted_idx.shape[0]) + .5
print(pos.size)
sorted_idx.size

```

```
plt.figure(figsize=(10,10))
plt.barh(pos, feature_importance[sorted_idx], align='center')
plt.yticks(pos, X.columns[sorted_idx])
plt.xlabel('Relative Importance')
plt.title('Variable Importance')
plt.show()
```

30



The popularity of random forest is primarily due to how well it performs in a multitude of data situations. It tends to handle highly correlated features well, where as a linear regression model would not. In this case study we demonstrate the performance ability even with only a few features and almost all of them being highly correlated with each other. Random Forest is also used as an efficient way to investigate the importance of a set of features with a large data set. Consider random forest to be one of your first choices when building a decision tree, especially for multiclass classifications.